Justin Bahr
NUID: 001592707

High Performance Computing
Homework 4 - Question 2

I ran my programs on the Explorer Cluster with 16 processes per node (except for the run with 32 bins on 4 nodes, which only uses 8 processes per node).

a. Below is an example result of program run with 128 bins:

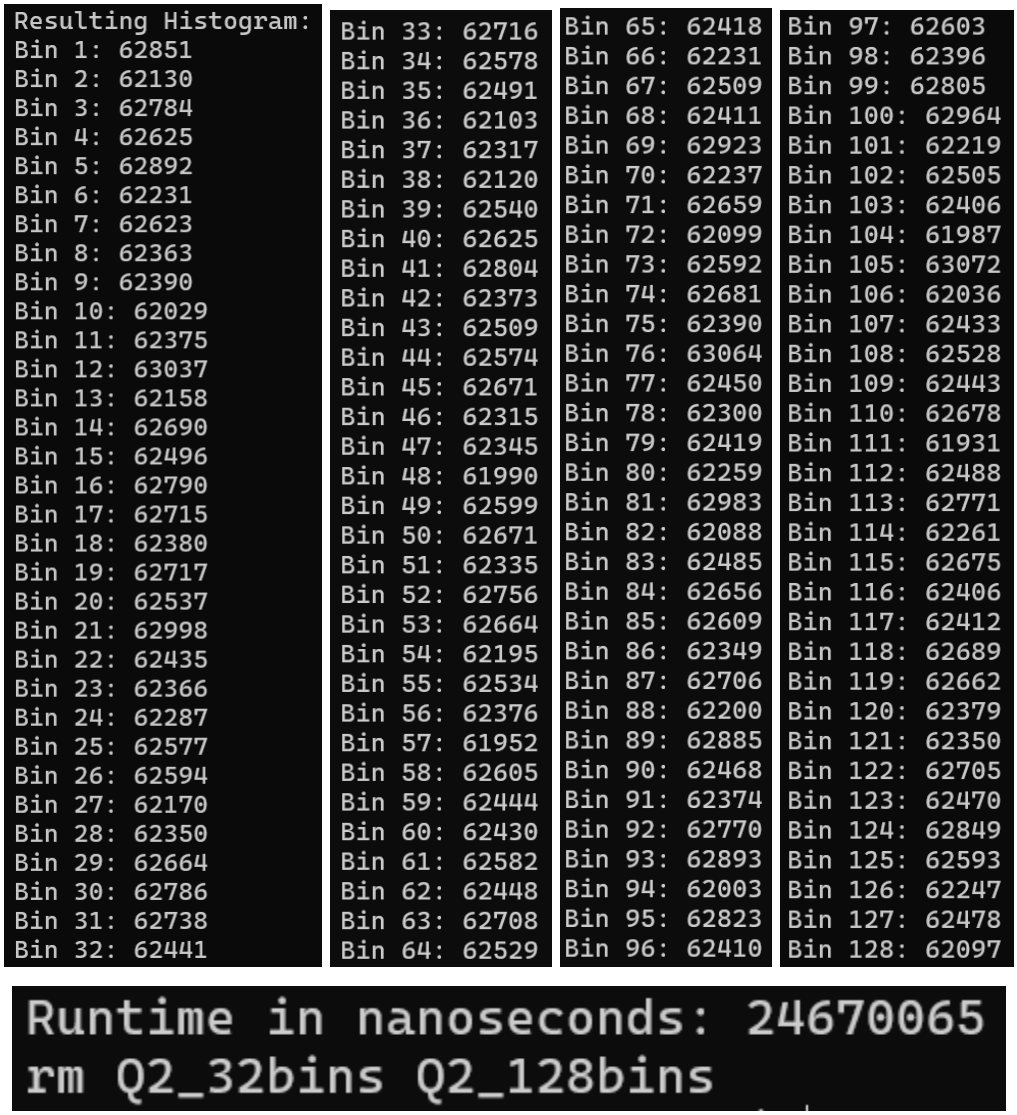

Figure 1. Screenshot of 128 Bins Result

Since there are 8 millions random integers to histogram and 128 bins, each bin should have about 8 million / 128 = 62,500 integers, which resembles the results above.

128 Bins	2 Nodes	4 Nodes	8 Nodes
Trial	Runtime (ms)	Runtime (ms)	Runtime (ms)
1	18.539344	24.811686	37.20247
2	18.650524	24.706874	27.847574
3	18.715954	24.656571	28.290499
Average	18.635274	24.72504367	31.11351433

Figure 2. Runtimes for 128 Bins with 2, 4, and 8 Nodes

In this experiment, increasing the number of nodes used also increased the runtime. While increasing the number of nodes used decreases the amount of work each node needs to do, it also increases the overhead cost of MPI communication between nodes, which is significantly longer than shared memory latency. However, this is using a fixed size of 8 million integers to judge strong scaling. If the problem size is increased past a point where the overhead communication costs are negligible, it will likely become advantageous to employ MPI across multiple nodes.

b. Below is an example result of program run with 32 bins

```
[bahr.j@c2189 Question2]$ cat slurm-107333.out
mpiCC -o Q2_32bins Q2_32bins.cpp
mpiCC -o Q2_128bins Q2_128bins.cpp

Resulting Histogram:
Bin 1: 250011
Bin 2: 250082
Bin 3: 249654
Bin 4: 249778
Bin 5: 250386
Bin 6: 249821
Bin 7: 250185
Bin 8: 249940
Bin 9: 250974
Bin 10: 249855
Bin 11: 250606
Bin 12: 249689
Bin 13: 249650
Bin 14: 250050
Bin 15: 249300
Bin 16: 249420
Bin 17: 250902
Bin 18: 249644
Bin 19: 249618
Bin 20: 251260
Bin 21: 250427
Bin 22: 249808
Bin 23: 250035
Bin 24: 249998
Bin 25: 249632
Bin 26: 249479
Bin 27: 249101
Bin 28: 249858
Bin 29: 250634
Bin 30: 250091
Bin 31: 250699
Bin 32: 249413

Runtime in nanoseconds: 16677251
rm Q2_32bins Q2_128bins
```

Figure 3. Screenshot of 32 Bins Result

Since there are 8 millions random integers to histogram and 32 bins, each bin should have about 8 million / 32 = 250,000 integers, which resembles the results above.

32 Bins	2 Nodes	4 Nodes
Trial	Runtime (ms)	Runtime (ms)
1	16.940292	23.131445
2	16.740868	24.088682
3	16.802406	23.62041
Average	16.82785533	23.61351233

Figure 2. Runtimes for 32 Bins with 2 and 4 Nodes

Similar to part (a), increasing the from 2 to 4 nodes also increased the runtime due to the overhead cost of MPI communication across nodes.

c. For both 2 and 4 nodes, histogramming ran faster with 32 bins compared to with 128 bins. This is because the local histogram arrays for each process only require 32 indices compared to 128. Increasing the size of the local histogram increases the size of the messages passed between MPI processes during histogram reduction, leading to longer runtimes. I confirmed this by increasing the number of bins to 256, which led to an increased runtime of 24.82 ms with 2 nodes compared to 18.64 ms with 2 nodes and 128 bins.